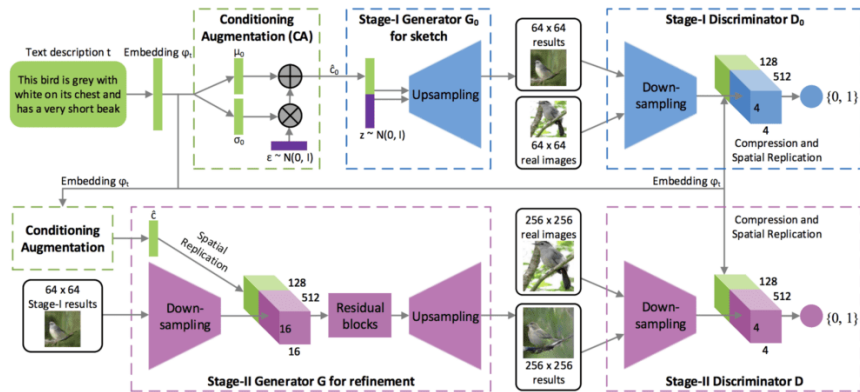


Deep Learning

ISLP Lab Solution



Kunal Gokhe
MT2309
2024

Contents

Single Layer Neural Network on Hitters Data	2
Multilayer Network on the MNIST Digit Data	7
Convolutional Neural Networks	13
Using Pretrained CNN Models	19
IMDB Document Classification	24
Recurrent Neural Networks	29
Time Series Prediction	34

Single Layer Neural Network on Hitters Data


Introduction

In this task, we will design the neural network for predicting the Salary based on the predictor variables. We will design the network with 22 input and single output with single hidden layers with 50 hidden units.

Dataset

The dataset consists of 23 columns and some of the columns contain categorical values, which we have to convert into the dummy variable using a one-hot encoding approach. No. of observations in the dataset are 263. The ratio of splitting dataset into training, validation and test are 80%, 10% and 10% respectively.

Neural Network Design



Code

Code for Importing Libraries

```
1 # Import Libraries
2 import numpy as np
3 import pandas as pd
```

```

4 import matplotlib.pyplot as plt
5 from matplotlib.animation import FuncAnimation
6 import torch
7 from torch.utils.data import DataLoader, random_split, TensorDataset
8 import torch.nn as nn
9 import torch.optim as optim
10 import torchvision
11 import torchvision.transforms as transforms
12 import torchinfo
13 from torchmetrics import R2Score, MeanAbsoluteError
14 from ISLP import load_data
15 from sklearn.model_selection import train_test_split

```

Selecting NVIDIA-CUDA Device to train on GPU

```

1 # Select Device for training
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3 print('Selected Device for Training: ', device)

```

Import Dataset for training model, we also standardize the dataset which will help NN to train better.

```

1 # Import Data
2 data=load_data('Hitters').dropna()
3 # Extract columns of category
4 catCols=data.select_dtypes(include='category').columns
5 # One-hot Encoding for Category data
6 data=pd.get_dummies(data,catCols)
7 # Compute the mean and standard deviation for each column
8 means = np.mean(data, axis=0)
9 stds = np.std(data, axis=0)
10 # Standardize the data
11 standardized_data = (data - means) / stds
12 # Split Response and Predictor variable
13 Xdata=standardized_data[standardized_data.drop(columns='Salary').columns].to_numpy()
14 Ydata=standardized_data['Salary'].to_numpy()

```

Now, Convert the dataset to the Tensor Dataloader which will be required for training the neural network model, Dataloader used to increase the performance of pre-processing of data if any and which pass the data in batch without sacrificing the performance of machine.

```

1 # Train, validate and test splits
2 Xtrain, Xtemp, Ytrain, Ytemp = train_test_split(Xdata, Ydata, test_size=0.2)
3 Xvalid, Xtest, Yvalid, Ytest = train_test_split(Xtemp, Ytemp, test_size=0.1/0.2)
4 # Create TensorDatasets
5 train_dataset = TensorDataset(torch.tensor(Xtrain, dtype=torch.float32), torch.
    tensor(Ytrain, dtype=torch.float32))
6 valid_dataset = TensorDataset(torch.tensor(Xvalid, dtype=torch.float32), torch.
    tensor(Yvalid, dtype=torch.float32))
7 test_dataset = TensorDataset(torch.tensor(Xtest, dtype=torch.float32), torch.
    tensor(Ytest, dtype=torch.float32))
8 # Create DataLoaders
9 train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
10 valid_loader = DataLoader(valid_dataset, batch_size=32, shuffle=False)
11 test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

```

Define Neural Network containing single hidden layer of 50 units

```

1 # NN Model
2 class Model(nn.Module):
3     def __init__(self):
4         super(Model,self).__init__()

```

```

5         self.flatten=nn.Flatten()
6         self.sequential=nn.Sequential(
7             nn.Linear(22,50),
8             nn.ReLU(),
9             nn.Dropout(0.4),
10            nn.Linear(50,1)
11        )
12    def forward(self,x):
13        x=self.flatten(x)
14        return torch.flatten(self.sequential(x))
15 # Transfer Model to CUDA Device
16 model = Model().to(device)

```

Set the Hyperparameters and performance metrics

```

1 # Hyperparameters and Performance Metrics
2 criterion = nn.MSELoss()
3 optimizer = optim.Adam(model.parameters(),lr=0.001)
4 r2_metric = R2Score().to(device)
5 mae_metric = MeanAbsoluteError().to(device)
6 Nepochs = 100

```

The training and validation section is shown below, it will plot the Loss for both on figure and update the figure on each epochs to see the progress of training.

```

1 # open Figure for Accuracies
2 fig,ax=plt.subplots()
3 ax.set_xlim(0,Nepochs)
4 ax.set_ylim(0,1)
5 ax.set_xlabel('Epochs')
6 ax.set_ylabel('MSE')
7 ax.set_title('Training and Validation Loss')
8 train_acc_line, = ax.plot([], [], label='Training')
9 valid_acc_line, = ax.plot([], [], label='Validation')
10 plt.legend()
11 plt.draw()
12
13 # Training the Model
14 trainLoss,validLoss=[],[]
15 for epoch in range(Nepochs):
16     model.train()
17     running_loss = 0
18     for input,target in train_loader:
19         optimizer.zero_grad()
20         # Add inputs and labels to device
21         input, target = input.to(device), target.to(device)
22         # Forward pass
23         predictions = model(input)
24         # Compute loss
25         loss = criterion(predictions,target)
26         # Backward pass
27         loss.backward()
28         optimizer.step()
29         # Compute loss
30         running_loss += loss.item()
31     trainLoss.append(running_loss/len(train_loader))
32     print(f'Epoch {epoch+1}/{Nepochs}, Loss: {running_loss/len(train_loader):.4f}')
33
34     model.eval()
35     running_loss = 0
36     with torch.no_grad():
37         for input,target in valid_loader:

```

```

38         # Add inputs and labels to device
39         input, target = input.to(device), target.to(device)
40         predictions = model(input)
41         # Compute loss
42         loss = criterion(predictions, target)
43         # Compute loss
44         running_loss += loss.item()
45         validLoss.append(running_loss/len(valid_loader))
46     # Update Plot
47     train_acc_line.set_data(range(1, len(trainLoss) + 1), trainLoss)
48     valid_acc_line.set_data(range(1, len(trainLoss) + 1), validLoss)
49     ax.relim()
50     plt.draw()
51     plt.pause(0.1)
52 plt.show()

```

Testing of the trained model is the essential part of the post-designing process, here we check how the model is performed on the test dataset.

```

1 # Test Model
2 model.eval()
3 testLoss=[]
4 running_loss = 0
5 with torch.no_grad():
6     for input,target in test_loader:
7         # Add inputs and labels to device
8         input, target = input.to(device), target.to(device)
9         predictions = model(input)
10        # Compute loss
11        loss = criterion(predictions, target)
12        r2_metric.update(predictions, target)
13        mae_metric.update(predictions, target)
14        # Compute loss
15        running_loss += loss.item()
16    testLoss=running_loss/len(test_loader)
17 print('Testing Results')
18 print(f'R-Squared: {r2_metric.compute():0.4f}')
19 print(f'MSE: {testLoss:0.4f}')
20 print(f'MAE: {mae_metric.compute():0.4f}')

```

Hyperparameters and Performance Metrics

Hyperparameters

Optimizer : ADAM

Epochs : 100

Criterion : MSE Loss

Batch Size : 32

Performance Metric

R-Squared : 0.7640

MSE : 0.3231

MAE : 0.4075



Conclusions

From the above results we can see that the model fit is good as the R-square values is close to 1 but the R-square value can be better, it is good if we have large amount of data to train model and get good prediction on salary. Relatively high value of R-square suggesting that the model fits the data well. It indicates that the neural network has successfully captured a substantial amount of the underlying patterns related to the performance of the hitters while predicting the Salary.

Multilayer Network on the MNIST Digit Data


Introduction


In this section we will create the multilayer network which will be trained on the MNIST digit dataset from 0 to 9. The image size of 28x28 is the input of the model and 10 output represent the probability of occurring that response number label. The model is trained on the GPU for faster computation and with batchsize of 128.

Dataset

The MNIST database is a large database of handwritten digits that is commonly used for training various image processing systems. The MNIST database of handwritten digits has a training set of 60,000 examples, and a test set of 10,000 examples. The train, valid and test set is splitted into 80%, 10% and 10% respectively.

Neural Network Design





Code

Import Libraries

```

1 # Import Libraries
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from matplotlib.animation import FuncAnimation
6 import torch
7 from torch.utils.data import DataLoader, random_split, TensorDataset
8 import torch.nn as nn
9 import torch.optim as optim
10 import torchvision
11 import torchvision.transforms as transforms
12 import torchinfo
13 from torchmetrics import R2Score, MeanAbsoluteError
14 from ISLP import load_data
15 from sklearn.model_selection import train_test_split

```

Selecting NVIDIA-CUDA Device to train on GPU

```

1 # Select Device for training
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3 print('Selected Device for Training: ', device)

```

Augment Data and Add to Dataloader

```

1 Data Augmentation Transform
2 transform = transforms.Compose([
3     transforms.ToTensor()
4 ])
5
6 # Load MNIST training and test datasets
7 train_dataset = torchvision.datasets.MNIST(root='./data', train=True, download=
    True, transform=transform)
8 train_dataset, valid_dataset = random_split(train_dataset, [int(0.8*len(train_dataset)
    ), int(0.2*len(train_dataset))])
9 test_dataset = torchvision.datasets.MNIST(root='./data', train=False, download=
    True, transform=transform)
10 # Data loaders for batching
11 train_loader = DataLoader(train_dataset, batch_size=128, shuffle=True)
12 valid_loader = DataLoader(valid_dataset, batch_size=128, shuffle=True)
13 test_loader = DataLoader(test_dataset, batch_size=128, shuffle=False)

```

Create Multilayer Neural Network to train MNIST Dataset

```
1 # NN Model
2 class Model(nn.Module):
3     def __init__(self):
4         super(Model, self).__init__()
5         self.layer1 = nn.Sequential(
6             nn.Flatten(),
7             nn.Linear(28*28, 256),
8             nn.ReLU(),
9             nn.Dropout(0.4))
10        self.layer2 = nn.Sequential(
11            nn.Linear(256, 128),
12            nn.ReLU(),
13            nn.Dropout(0.3))
14        self._forward = nn.Sequential(
15            self.layer1,
16            self.layer2,
17            nn.Linear(128, 10))
18    def forward(self, x):
19        return self._forward(x)
20
21 # Transfer Model to CUDA Device
22 model = Model().to(device)
23 torchinfo.summary(model)
```

Set the Hyperparameters and performance metrics

```
1 # Hyperparameters and Performance Metrics
2 criterion = nn.CrossEntropyLoss()
3 optimizer = optim.Adam(model.parameters(), lr=0.001)
4 Nepochs = 10
```

The training and validation section is shown below, it will plot the Loss for both on figure and update the figure on each epochs to see the progress of training.

```
1 # open Figure for Accuracies
2 fig,ax=plt.subplots()
3 ax.set_xlim(0,Nepochs)
4 ax.set_ylim(0,100)
5 ax.set_xlabel('Epochs')
6 ax.set_ylabel('Accuracy (%)')
7 ax.set_title('Training and Validation Accuracy')
8 train_acc_line, = ax.plot([], [], label='Training')
9 valid_acc_line, = ax.plot([], [], label='Validation')
10 plt.legend()
11 plt.draw()
12
13 # Training the Model
14 trainAcc,validAcc=[],[]
15 for epoch in range(Nepochs):
16     model.train()
17     correct_train = 0
18     total_train = 0
19     running_loss = 0
20     for input,target in train_loader:
21         optimizer.zero_grad()
22         # Add inputs and labels to device
23         input, target = input.to(device), target.to(device)
24         # Forward pass
25         predictions = model(input)
26         # Compute loss
27         loss = criterion(predictions,target)
```

```

28         # Backward pass
29         loss.backward()
30         optimizer.step()
31         # Compute loss
32         running_loss += loss.item()
33         # Compute Accuracy
34         _, predicted = torch.max(predictions.data, 1)
35         total_train += target.size(0)
36         correct_train += (predicted == target).sum().item()
37     # Store Training Accuracy
38     trainAcc.append(correct_train / total_train * 100)
39     print(f'Epoch {epoch+1}/{Nepochs}, Loss: {running_loss/len(train_loader):.4f}'
40         )
41
42     model.eval()
43     running_loss = 0
44     correct_valid = 0
45     total_valid = 0
46     with torch.no_grad():
47         for input,target in valid_loader:
48             # Add inputs and labels to device
49             input, target = input.to(device), target.to(device)
50             predictions = model(input)
51             # Compute loss
52             loss = criterion(predictions, target)
53             # Compute loss
54             running_loss += loss.item()
55             # Compute Accuracy
56             _, predicted = torch.max(predictions.data, 1)
57             total_valid += target.size(0)
58             correct_valid += (predicted == target).sum().item()
59         validAcc.append(correct_valid / total_valid * 100)
60     # Update Plot
61     train_acc_line.set_data(range(1, len(trainAcc) + 1), trainAcc)
62     valid_acc_line.set_data(range(1, len(trainAcc) + 1), validAcc)
63     ax.relim()
64     plt.draw()
65     plt.pause(0.1)
66     plt.show()
67     print('Training Accuracy: ',trainAcc[-1])

```

Testing of the trained model is the essential part of the post-designing process, here we check how the model is performed on the test dataset.

```

1 # Test Model
2 model.eval()
3 testLoss=[]
4 correct_test = 0
5 total_test = 0
6 running_loss = 0
7 with torch.no_grad():
8     for input,target in test_loader:
9         # Add inputs and labels to device
10        input, target = input.to(device), target.to(device)
11        predictions = model(input)
12        # Compute loss
13        loss = criterion(predictions, target)
14        # Compute loss
15        running_loss += loss.item()
16        total_test += target.size(0)
17        _, predicted = torch.max(predictions.data, 1)
18        correct_test += (predicted == target).sum().item()

```

```

19     testLoss=running_loss/len(test_loader)
20     print('Testing Accuracy: ', correct_test / total_test * 100)

```

We will also show the random image of data and plot the actual and predicted labels as shown below.

```

1  # Move images back to CPU for plotting
2  images = input.cpu().numpy()
3  # Plot the images and their predictions
4  fig = plt.figure(figsize=(12, 12))
5  for i in range(9):
6      plt.subplot(3, 3, i+1)
7      plt.imshow(np.squeeze(images[i]), cmap='gray')
8      plt.title(f"Pred: {predicted[i].item()}, Actual: {target[i].item()}")
9      plt.axis('off')
10 plt.show()

```

Hyperparameters and Performance Metrics

Hyperparameters

Optimizer : ADAM

Epochs : 10


Criterion : Cross Entropy Loss


Batch Size : 128

Performance Metric

Training Accuracy : 97.80

Testing Accuracy : 97.86





Conclusions

From the above results, we can say that the model is performing very well in predicting the number from the given image, the accuracy of the model is 97% in just 10 epochs which gives us good results. The testing accuracy of the model is also promising which is 97.86%. So, we can concluded that the designed NN model gives the good results for the given MNIST dataset.

Convolutional Neural Networks


Introduction


In this part we will design the CNN Model to classify the collection of the natural image from the dataset, we will use CIFAR100 dataset which contains 100 class of similar kind of image, we will design CNN model which takes the input image and at output we get the predicted class.

Dataset

The dataset consists of 50,000 training images and 10,000 testing image. out of the training we have use 10,000 for validating the model. The input image is first augmented using the transform by giving the random horizontal and vertical flips after that the Dataloader will handle the image with batchsize of 128.

Neural Network Design





Code

Import Libraries

```

1 # Import Libraries
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from matplotlib.animation import FuncAnimation
6 import torch
7 from torch.utils.data import DataLoader, random_split
8 import torch.nn as nn
9 import torch.optim as optim
10 import torchvision
11 import torchvision.transforms as transforms
12 import torchinfo

```

Selecting NVIDIA-CUDA Device to train on GPU

```

1 # Select Device for training
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3 print('Selected Device for Training: ', device)

```

Data Augmentation and Create Dataloader

```

1 # Data Augmentation Transform

```

```

2 transform = transforms.Compose([
3     transforms.RandomHorizontalFlip(),
4     transforms.ToTensor()
5 ])
6 # Import Dataset
7 train_dataset = torchvision.datasets.CIFAR100(root='./data', train=True, download=
    True, transform=transform)
8 train_dataset,valid_dataset=random_split(train_dataset,[int(0.8*len(train_dataset)
    ),int(0.2*len(train_dataset))])
9 test_dataset = torchvision.datasets.CIFAR100(root='./data', train=False, download=
    True, transform=transform)
10 # Create DataLoader
11 train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=128, shuffle=
    True)
12 valid_loader = torch.utils.data.DataLoader(valid_dataset, batch_size=128, shuffle=
    True)
13 test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=128, shuffle=
    False)

```

Define Neural Network containing the CNN Model

```

1 # Create NN Model
2 class BuildingBlock(nn.Module):
3     def __init__(self,in_channels,out_channels):
4         super(BuildingBlock , self).__init__()
5         self.conv = nn.Conv2d(in_channels=in_channels,out_channels=out_channels ,
6             kernel_size =(3 ,3),padding='same')
7         self.activation = nn.ReLU ()
8         self.pool = nn.MaxPool2d(kernel_size =(2 ,2))
9         def forward(self , x):
10             return self.pool(self.activation(self.conv(x)))
11 class CIFARModel(nn.Module):
12     def __init__(self):
13         super(CIFARModel , self).__init__()
14         sizes = [(3 ,32) ,
15             (32 ,64) ,
16             (64 ,128) ,
17             (128 ,256)]
18         self.conv = nn.Sequential (*[ BuildingBlock(in_ , out_) for in_ , out_ in
19             sizes ])
20         self.output = nn.Sequential(nn.Dropout (0.5) ,
21             nn.Linear (2*2*256 , 512) ,
22             nn.ReLU () ,
23             nn.Linear (512, 100))
24         def forward(self , x):
25             val = self.conv(x)
26             val = torch.flatten(val , start_dim =1)
27             return self.output(val)
28 # Model Summary
29 model = CIFARModel().to(device)
30 torchinfo.summary(model,input_size=(128,3,32,32),col_names=['input_size','
    output_size','num_params'])

```

Set the Hyperparameters and performance metrics

```

1 criterion = nn.CrossEntropyLoss()
2 optimizer = optim.Adam(model.parameters(), lr=0.001)
3 Nepochs=50

```

The training and validation section is shown below, it will plot the Accuracy for both on figure and update the figure on each epochs to see the progress of training.


```

1  # open Figure for Accuracies
2  fig,ax=plt.subplots()
3  ax.set_xlim(0,Nepochs)
4  ax.set_ylim(0,100)
5  ax.set_xlabel('Epochs')
6  ax.set_ylabel('Accuracy (%)')
7  ax.set_title('Training and Validation Accuracy')
8  train_acc_line, = ax.plot([], [], label='Training')
9  valid_acc_line, = ax.plot([], [], label='Validation')
10 plt.legend()
11 plt.draw()
12
13 # Train Model
14 trainAcc,validAcc=[],[]
15 for epoch in range(Nepochs):
16     model.train()
17     correct_train = 0
18     total_train = 0
19     running_loss = 0
20     for i, (inputs, labels) in enumerate(train_loader):
21         # Add inputs and labels to device
22         inputs, labels = inputs.to(device), labels.to(device)
23         # Zero the parameter gradients
24         optimizer.zero_grad()
25         # Forward pass
26         outputs = model(inputs)
27         # Compute Loss
28         loss = criterion(outputs, labels)
29         # Backward pass
30         loss.backward()
31         optimizer.step()
32         # Compute Accuracy
33         _, predicted = torch.max(outputs.data, 1)
34         total_train += labels.size(0)
35         correct_train += (predicted == labels).sum().item()
36
37         # Compute Batch Loss
38         running_loss += loss.item()
39         # Print every 100 mini-batches
40         if i % 100 == 99:
41             print(f'Epoch [{epoch+1}/{Nepochs}], Step [{i+1}/{len(train_loader)}], Loss: {running_loss/100:.4f}')
42             running_loss = 0
43         # Store Training Accuracy
44         trainAcc.append(correct_train / total_train * 100)
45
46     # Validate Model
47     model.eval()
48     correct_valid = 0
49     total_valid = 0
50     with torch.no_grad():
51         for inputs, labels in valid_loader:
52             # Add inputs and labels to device
53             inputs, labels = inputs.to(device), labels.to(device)
54             # Forward pass
55             outputs = model(inputs)
56             # Compute Accuracy
57             _, predicted = torch.max(outputs.data, 1)
58             total_valid += labels.size(0)
59             correct_valid += (predicted == labels).sum().item()
60         # Store Validation Accuracy
61         validAcc.append(correct_valid / total_valid * 100)

```

```

62
63     # Update Plot
64     train_acc_line.set_data(range(1, len(trainAcc)+1), trainAcc)
65     valid_acc_line.set_data(range(1, len(trainAcc) + 1), validAcc)
66     plt.draw()
67     plt.pause(0.1)
68 plt.show()
69 print('Training Accuracy: ', trainAcc[-1])

```

Testing of the trained model is the essential part of the post-designing process, here we check how the model is performed on the test dataset.

```

1 # Test Model
2 model.eval()
3 correct_test = 0
4 total_test = 0
5 with torch.no_grad():
6     for inputs, labels in test_loader:
7         # Add inputs and labels to device
8         inputs, labels = inputs.to(device), labels.to(device)
9         # Forward pass
10        outputs = model(inputs)
11        # Compute Accuracy
12        _, predicted = torch.max(outputs.data, 1)
13        total_test += labels.size(0)
14        correct_test += (predicted == labels).sum().item()
15 print('Testing Accuracy: ', correct_test / total_test * 100)

```

Hyperparameters and Performance Metrics

Hyperparameters

Optimizer : ADAM

Epochs : 50



Criterion : Cross Entropy Loss

Batch Size : 128

Performance Metric

Training Accuracy : 61.21

Testing Accuracy : 45.68



Conclusions

We concluded that using the developed CNN Model we got a decent accuracy of 67.21% but the testing accuracy is not good enough to use the model, the reason for the poor performance may be due to lesser data size for each classes in dataset or we can apply some more dense layer to get the good results. From the tested results we can see that some images get easily able to predict but it is also tricky to predict from some classes where the feature is very similar.

Using Pretrained CNN Models


Introduction

In this task we are going to use the pre-trained model ResNet-50 to predict the class of the given image. we import the model with trained weights for the imagenet dataset. using the few images we will test the model and check how the ResNet-50 Model performs.

Dataset

Here we are not going to train the model but as far the dataset is concerned, ResNet-50 model is trained on the ImageNet dataset consisting of 1000 classes and 14 Million images. We are going to use few image to test the Model response.

Neural Network Design



Code

Import Libraries

```
1 # Import Libraries
2 import numpy as np
3 import pandas as pd
4 import os
5 import matplotlib.pyplot as plt
6 from matplotlib.animation import FuncAnimation
7 import torch
8 from torch.utils.data import DataLoader, random_split, Dataset
9 import torch.nn as nn
10 import torch.optim as optim
11 import torchvision
12 import torchvision.transforms as transforms
13 import torchinfo
14 from PIL import Image
```

Selecting NVIDIA-CUDA Device to train on GPU

```
1 # Select Device for training
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3 print('Selected Device for Training: ',device)
```

Data Augmentation and DataLoader

```

1 # Data Augmentation Transform
2 transform = transforms.Compose([
3     transforms.Resize((232,232)),
4     transforms.CenterCrop(224),
5     transforms.ToTensor(),
6     # transforms.Normalize([0.485 ,0.456 ,0.406] ,[0.229 ,0.224 ,0.225])
7 ])
8 # Image Dataset Loader
9 class ImageDataset(Dataset):
10     def __init__(self, image_folder, transform=None):
11         self.image_folder = image_folder
12         self.image_files = [f for f in os.listdir(image_folder) if f.endswith(('
13         jpg','png'))]
14         self.transform = transform
15
16     def __len__(self):
17         return len(self.image_files)
18
19     def __getitem__(self, idx):
20         img_path = os.path.join(self.image_folder, self.image_files[idx])
21         image = Image.open(img_path) # Open image
22         if self.transform:
23             image = self.transform(image) # Apply transformation
24         return image
25
26 # Import Dataset
27 test_dataset = ImageDataset('./data/book_images',transform=transform)
28 # Create DataLoader
29 test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=1, shuffle=
30     False)

```

ResNet-50 Model

```

1 # Resnet50 Model
2 model=torchvision.models.resnet50(pretrained=True).to(device)
3 # Extract Categories of ImageNet
4 category=torchvision.models.ResNet50_Weights.DEFAULT.meta['categories']

```




Evaluate model

```


1 # Test Model
2 model.eval()
3 for i,inputs in enumerate(test_loader):
4     # Add inputs and labels to device
5     inputs = inputs.to(device)
6     # Forward pass
7     outputs = model(inputs)
8     prediction = outputs.squeeze(0).softmax(0)
9     classid = prediction.argmax().item()
10    predicted_label = category[classid]
11    # Find 5 strongest prediction
12    nclassid=prediction.argsort(descending=True)[0:5].cpu().numpy()
13    npredict_label= [category[i] for i in nclassid]
14    npred=[prediction[i].item() for i in nclassid]
15    img = torchvision.utils.make_grid(inputs.cpu()).numpy()
16    fig, axes = plt.subplots(1, 2, figsize=(12,8))
17    axes[0].imshow(np.transpose(img, (1, 2, 0)))
18    axes[0].set_title(f'Predicted: {predicted_label}', fontsize=10)
19    axes[0].axis('off')
20    axes[1].bar(npredict_label,npred)
21    axes[1].tick_params(axis='x',rotation = 45)
22    plt.show()

```


Results





Predicted: kite



Predicted: kite





Conclusions

We have predict the class of the image using the ResNet-50 pre-trained network, we can see that it model is able to works well for the given inputs which is used for testing. in the first image instead of cape weaver it is predicted as cricket, flamingo image is predicted correctly, In 3rd and 4th image it predicts the wrong which should be Cooper's hawk. So, the ResNet-50 done the reasonable job for classifying the given image.

IMDB Document Classification


Introduction

In this we will create the NN model which can predict the sentiment based on the given text review of movie. for this task we are going to use IMDB dataset which has review of movie with sentiment of review. In this we use tokenization of word and Vocabulary library to embed the text data for training the model.

Dataset

We have a training set and test set, each with 25,000 examples, and each balanced with regard to sentiment. The resulting training feature matrix X has dimension 25,000 10,000, but only 1.3% of the binary entries are non zero. We split off a validation set in ratio of 10% from the 25,000 training observations.

Neural Network Design



Code

Import Libraries

```
1 # Import Libraries
```

```

2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from matplotlib.animation import FuncAnimation
6 import torch
7 from torch.utils.data import DataLoader, random_split
8 import torch.nn as nn
9 import torch.optim as optim
10 import torchvision
11 import torchvision.transforms as transforms
12 from torchtext import data, datasets # Use legacy for compatibility
13 import torchinfo
14 import spacy

```

Selecting NVIDIA-CUDA Device to train on GPU

```

1 # Select Device for training
2 device = torch.device('cpu')
3 print('Selected Device for Training: ', device)

```

Tokenizing Words

```

1 # Load the SpaCy English tokenizer
2 spacy_en = spacy.load('en_core_web_sm')
3 # Define a function to tokenize the text using SpaCy
4 def tokenize_en(text):
5     return [tok.text for tok in spacy_en.tokenizer(text)]

```

Create DataLoader for Text Dataset

```

1 # Define fields for text and labels
2 TEXT = data.Field(tokenize=tokenize_en, lower=True)
3 LABEL = data.LabelField(dtype=torch.float)
4 # Load IMDB dataset
5 train_data, test_data = datasets.IMDB.splits(TEXT, LABEL)
6 train_data, valid_data = train_data.split(split_ratio=0.8)
7 # Build the vocabulary for the text and labels
8 TEXT.build_vocab(train_data, max_size=10000, vectors="glove.6B.100d", unk_init=
    torch.Tensor.normal_)
9 LABEL.build_vocab(train_data)
10 # Create an iterator for training data
11 train_iterator, valid_iterator, test_iterator = data.BucketIterator.splits((
    train_data, valid_data, test_data), batch_size=64, device=device)

```

Define Neural Network

```

1 # NN Model
2 class IMDBModel(nn.Module):
3     def __init__(self, input_size):
4         super(IMDBModel, self).__init__()
5         self.dense1 = nn.Linear(input_size, 16)
6         self.activation = nn.ReLU()
7         self.dense2 = nn.Linear(16, 16)
8         self.output = nn.Linear(16, 1)
9
10    def forward(self, x):
11        val = x.mean(dim=0)
12        for _map in [self.dense1, self.activation, self.dense2, self.activation,
13            self.output]:
14            val = _map(val)
15        return torch.flatten(val)
16
17 model = IMDBModel(100) # Input size of 100 (as we are using GloVe 100d vectors)
18 model = model.to(device)

```

Set the Hyperparameters and performance metrics

```
1 # Hyperparameters and Performance Metrics
2 criterion = nn.BCEWithLogitsLoss()
3 optimizer = optim.Adam(model.parameters(), lr=0.001)
4 Nepochs = 50
```

The training and validation section is shown below, it will plot the Accuracy for both on figure and update the figure on each epochs to see the progress of training.

```
1 # open Figure for Accuracies
2 fig, ax = plt.subplots()
3 ax.set_xlim(0, Nepochs)
4 ax.set_ylim(0, 100)
5 ax.set_xlabel('Epochs')
6 ax.set_ylabel('Accuracy (%)')
7 ax.set_title('Training and Validation Accuracy')
8 train_acc_line, = ax.plot([], [], label='Training')
9 valid_acc_line, = ax.plot([], [], label='Validation')
10 plt.legend()
11 plt.draw()
12
13 # Training the Model
14 trainAcc, validAcc = [], []
15 for epoch in range(Nepochs):
16     model.train()
17     correct_train = 0
18     total_train = 0
19     running_loss = 0
20     for batch in train_iterator:
21         optimizer.zero_grad()
22         # Get inputs and labels
23         text = batch.text
24         text = text.to(device)
25         labels = batch.label.to(device)
26         # Get embeddings from GloVe vectors
27         embedded = TEXT.vocab.vectors[text]
28         embedded = embedded.to(device)
29
30         # Forward pass
31         predictions = model(embedded)
32
33         # Compute loss
34         loss = criterion(predictions, labels)
35
36         # Backward pass
37         loss.backward()
38         optimizer.step()
39
40         # Compute loss
41         running_loss += loss.item()
42
43         # Compute accuracy
44         sigmoid_predictions = torch.sigmoid(predictions)
45         binary_predictions = (sigmoid_predictions >= 0.5).float()
46
47         correct_train += (binary_predictions == labels).sum().item()
48         total_train += labels.size(0)
49     trainAcc.append(correct_train / total_train * 100)
50     print(f'Epoch {epoch+1}/{Nepochs}, Loss: {running_loss:.4f}')
51
52     model.eval()
53     correct_valid = 0
```

```

54     total_valid = 0
55     with torch.no_grad():
56         for batch in valid_iterator:
57             text = batch.text
58             text = text.to(device)
59             labels = batch.label.to(device)
60
61             embedded = TEXT.vocab.vectors[text]
62             embedded = embedded.to(device)
63
64             predictions = model(embedded)
65
66             # Compute accuracy
67             sigmoid_predictions = torch.sigmoid(predictions)
68             binary_predictions = (sigmoid_predictions >= 0.5).float()
69             correct_valid += (binary_predictions == labels).sum().item()
70             total_valid += labels.size(0)
71         validAcc.append(correct_valid / total_valid * 100)
72
73     # Update Plot
74     train_acc_line.set_data(range(1, len(trainAcc) + 1), trainAcc)
75     valid_acc_line.set_data(range(1, len(trainAcc) + 1), validAcc)
76     plt.draw()
77     plt.pause(0.1)
78 plt.show()
79 print('Training Accuracy: ', trainAcc[-1])

```

Testing of the trained model is the essential part of the post-designing process, here we check how the model is performed on the test dataset.

```

1 # Test Model
2 model.eval()
3 correct_test = 0
4 total_test = 0
5 with torch.no_grad():
6     for batch in test_iterator:
7         text = batch.text
8         text = text.to(device)
9         labels = batch.label.to(device)
10
11         embedded = TEXT.vocab.vectors[text]
12         embedded = embedded.to(device)
13
14         predictions = model(embedded)
15
16         # Compute accuracy
17         sigmoid_predictions = torch.sigmoid(predictions)
18         binary_predictions = (sigmoid_predictions >= 0.5).float()
19         correct_test += (binary_predictions == labels).sum().item()
20         total_test += labels.size(0)
21 print('Testing Accuracy: ', correct_test / total_test * 100)

```

Hyperparameters and Performance Metrics

Hyperparameters

Optimizer : ADAM

Epochs : 50


Criterion : BCE With Logits Loss

Batch Size : 64

Performance Metric

Training Accuracy : 71.505

Testing Accuracy : 72.48



Conclusions

From the above results, we can say that model is working sufficiently well to predict that the review of the movie is positive or negative, the training accuracy is 71% and testing accuracy is 72% for just 50 epochs which is good enough for the prediction, more number of data can be incorporated to get the better prediction performance.

Recurrent Neural Networks


Introduction

In this section we will use the same IMDB dataset and train it using LSTM network and we will compare the performance of both the NN model on same dataset. The last model give 71% accuracy for 50 Epochs.

Dataset

We have a training set and test set, each with 25,000 examples, and each balanced with regard to sentiment. The resulting training feature matrix X has dimension 25,000 10,000, but only 1.3% of the binary entries are non zero. We split off a validation set in ratio of 10% from the 25,000 training observations.

Neural Network Design



Code

Import Libraries

```
1 # Import Libraries
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from matplotlib.animation import FuncAnimation
6 import torch
7 from torch.utils.data import DataLoader, random_split
8 import torch.nn as nn
9 import torch.optim as optim
```

```

10 import torchvision
11 import torchvision.transforms as transforms
12 from torchtext import data, datasets
13 import torchinfo
14 import spacy

```

Selecting NVIDIA-CUDA Device to train on GPU

```

1 # Select Device for training
2 device = torch.device('cpu')
3 print('Selected Device for Training: ', device)

```

Tokenizing Words and create vocabulary

```

1 # Load the SpaCy English tokenizer
2 spacy_en = spacy.load('en_core_web_sm')
3 # Define a function to tokenize the text using SpaCy
4 def tokenize_en(text):
5     return [tok.text for tok in spacy_en.tokenizer(text)]
6 # Define fields for text and labels
7 TEXT = data.Field(tokenize=tokenize_en, lower=True)
8 LABEL = data.LabelField(dtype=torch.float)
9 # Load IMDB dataset
10 train_data, test_data = datasets.IMDB.splits(TEXT, LABEL)
11 train_data, valid_data = train_data.split(split_ratio=0.8)
12 # Build the vocabulary for the text and labels
13 TEXT.build_vocab(train_data, max_size=10000, vectors="glove.6B.100d", unk_init=
    torch.Tensor.normal_)
14 LABEL.build_vocab(train_data)
15 # Create an iterator for training data
16 train_iterator, valid_iterator, test_iterator = data.BucketIterator.splits((
    train_data, valid_data, test_data), batch_size=64, device=device)

```

Create NN Model using LSTM Architecture

```

1 # NN Model
2 class LSTMModel(nn.Module):
3     def __init__(self, input_size):
4         super(LSTMModel, self).__init__()
5         self.embedding = nn.Embedding(input_size, 32)
6         self.lstm = nn.LSTM(input_size=100,
7                             hidden_size=32,
8                             batch_first=True)
9         self.dense = nn.Linear(32, 1)
10    def forward(self, x):
11        val, (h_n, c_n) = self.lstm(x.mean(dim=0))
12        return torch.flatten(self.dense(val))
13 model = LSTMModel(100) # Input size of 100 (as we are using GloVe 100d vectors)
14 model = model.to(device)

```

Set the Hyperparameters and performance metrics

```

1 # Hyperparameter and Performance Metrics
2 criterion = nn.BCEWithLogitsLoss()
3 optimizer = optim.Adam(model.parameters(), lr=0.001)
4 Nepochs = 50

```

The training and validation section is shown below, it will plot the Accuracy for both on figure and update the figure on each epochs to see the progress of training.

```

1 # open Figure for Accuracies
2 fig, ax = plt.subplots()
3 ax.set_xlim(0, Nepochs)
4 ax.set_ylim(0, 100)

```

```

5 ax.set_xlabel('Epochs')
6 ax.set_ylabel('Accuracy (%)')
7 ax.set_title('Training and Validation Accuracy')
8 train_acc_line, = ax.plot([], [], label='Training')
9 valid_acc_line, = ax.plot([], [], label='Validation')
10 plt.legend()
11 plt.draw()
12
13 # Training the Model
14 trainAcc, validAcc = [], []
15 for epoch in range(Nepochs):
16     model.train()
17     correct_train = 0
18     total_train = 0
19     running_loss = 0
20     for batch in train_iterator:
21         optimizer.zero_grad()
22         # Get inputs and labels
23         text = batch.text
24         text = text.to(device)
25         labels = batch.label.to(device)
26         # Get embeddings from GloVe vectors
27         embedded = TEXT.vocab.vectors[text]
28         embedded = embedded.to(device)
29
30         # Forward pass
31         predictions = model(embedded)
32
33         # Compute loss
34         loss = criterion(predictions, labels)
35
36         # Backward pass
37         loss.backward()
38         optimizer.step()
39
40         # Compute loss
41         running_loss += loss.item()
42
43         # Compute accuracy
44         sigmoid_predictions = torch.sigmoid(predictions)
45         binary_predictions = (sigmoid_predictions >= 0.5).float()
46
47         correct_train += (binary_predictions == labels).sum().item()
48         total_train += labels.size(0)
49     trainAcc.append(correct_train / total_train*100)
50     print(f'Epoch {epoch+1}/{Nepochs}, Loss: {running_loss:.4f}')
51
52     model.eval()
53     correct_valid = 0
54     total_valid = 0
55     with torch.no_grad():
56         for batch in valid_iterator:
57             text = batch.text
58             text = text.to(device)
59             labels = batch.label.to(device)
60
61             embedded = TEXT.vocab.vectors[text]
62             embedded = embedded.to(device)
63
64             predictions = model(embedded)
65
66             # Compute accuracy

```



```

67         sigmoid_predictions = torch.sigmoid(predictions)
68         binary_predictions = (sigmoid_predictions >= 0.5).float()
69         correct_valid += (binary_predictions == labels).sum().item()
70         total_valid += labels.size(0)
71         validAcc.append(correct_valid / total_valid * 100)
72
73     # Update Plot
74     train_acc_line.set_data(range(1, len(trainAcc) + 1), trainAcc)
75     valid_acc_line.set_data(range(1, len(trainAcc) + 1), validAcc)
76     plt.draw()
77     plt.pause(0.1)
78 plt.show()
79 print('Training Accuracy: ', trainAcc[-1])

```

Testing of the trained model is the essential part of the post-designing process, here we check how the model is performed on the test dataset.

```

1  # Test Model
2  model.eval()
3  correct_test = 0
4  total_test = 0
5  with torch.no_grad():
6      for batch in test_iterator:
7          text = batch.text
8          text = text.to(device)
9          labels = batch.label.to(device)
10
11         embedded = TEXT.vocab.vectors[text]
12         embedded = embedded.to(device)
13
14         predictions = model(embedded)
15
16         # Compute accuracy
17         sigmoid_predictions = torch.sigmoid(predictions)
18         binary_predictions = (sigmoid_predictions >= 0.5).float()
19         correct_test += (binary_predictions == labels).sum().item()
20         total_test += labels.size(0)
21 print('Testing Accuracy: ', correct_test / total_test * 100)

```

Hyperparameters and Performance Metrics

Hyperparameters

Optimizer : ADAM

Epochs : 50


Criterion : BCE With Logits Loss

Batch Size : 64

Performance Metric

Training Accuracy : 76.645

Testing Accuracy : 76.896



Conclusions

We can observed that there is increase in performance of 5% as compare to the previous method which use embedding. We still have a somewhat "entry level" RNN in spite of this additional LSTM complexity. With a different model size, different regularization, and more hidden layers, we could probably get marginally better results. Nevertheless, parameter optimization and exploring various architectures are laborious due to the lengthy training times of LSTM models.

Time Series Prediction

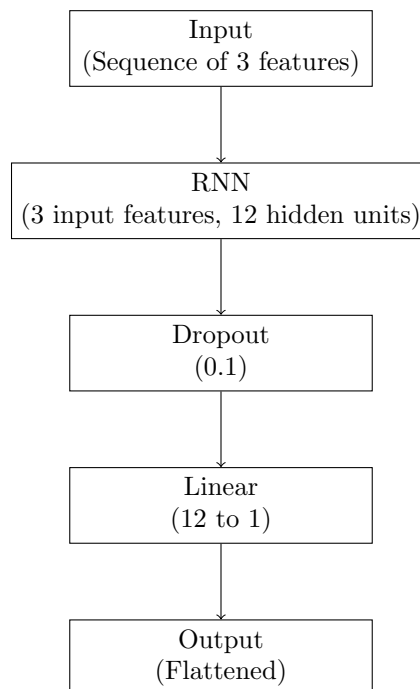
Introduction

Time series forecasting involves predicting future values of a given series based on its historical values. In this part we will take the NYSE dataset with three features and 5 lagged sequence and try to forecast the log volume of the NYSE stock.

Dataset

Historical trading statistics from the New York Stock Exchange. Daily values of the normalized log trading volume, DJIA return, and log volatility are given for a 24-year period from 1962–1986. The training, validation, and testing set are split in 80:10:10 ratio respectively.

Neural Network Design



Code

Import Libraries

```
1 # Import Libraries
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from matplotlib.animation import FuncAnimation
6 import torch
7 from torch.utils.data import DataLoader, random_split, TensorDataset
8 import torch.nn as nn
9 import torch.optim as optim
10 import torchvision
```

```

11 import torchvision.transforms as transforms
12 from torchtext import data, datasets
13 import torchinfo
14 import spacy
15 from ISLP import load_data
16 from sklearn.model_selection import train_test_split

```

Selecting NVIDIA-CUDA Device to train on GPU

```

1 # Select Device for training
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3 print('Selected Device for Training: ', device)

```

Import and Standardize dataset

```

1 # Import Dataset
2 NYSE = load_data('NYSE').dropna()
3 cols = ['DJ_return', 'log_volume', 'log_volatility']
4 data = NYSE[cols].values
5 # Compute the mean and standard deviation for each column
6 means = np.mean(data, axis=0)
7 stds = np.std(data, axis=0)
8 # Standardize the data
9 standardized_data = (data - means) / stds

```

Create lagged Feature and Splits the datasets

```

1 # Function to create lagged feature
2 def create_rnn_input(data, num_lags=5):
3     N = data.shape[0] - num_lags
4     rnn_input = np.zeros((N, num_lags, data.shape[1]))
5     next_volume = np.zeros(N)
6     for i in range(N):
7         rnn_input[i] = data[i:i + num_lags]
8         next_volume[i] = data[i + num_lags, 1]
9     return rnn_input, next_volume
10
11 # Train, valid and test split
12 Xdata, Ydata = create_rnn_input(standardized_data, num_lags=5)
13 Xtrain, Xtemp, Ytrain, Ytemp = train_test_split(Xdata, Ydata, test_size=0.2)
14 Xvalid, Xtest, Yvalid, Ytest = train_test_split(Xtemp, Ytemp, test_size=0.1/0.2)

```

Create DataLoader

```

1 # Create TensorDatasets
2 train_dataset = TensorDataset(torch.tensor(Xtrain, dtype=torch.float32), torch.
    tensor(Ytrain, dtype=torch.float32))
3 valid_dataset = TensorDataset(torch.tensor(Xvalid, dtype=torch.float32), torch.
    tensor(Yvalid, dtype=torch.float32))
4 test_dataset = TensorDataset(torch.tensor(Xtest, dtype=torch.float32), torch.
    tensor(Ytest, dtype=torch.float32))
5 # Create DataLoaders
6 train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
7 valid_loader = DataLoader(valid_dataset, batch_size=64, shuffle=False)
8 test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)

```

Define NN Model using RNN

```

1 # NN Model
2 class NYSEModel(nn.Module):
3     def __init__(self):
4         super(NYSEModel, self).__init__()
5         self.rnn = nn.RNN(3, 12, batch_first=True)
6         self.dense = nn.Linear(12, 1)
7         self.dropout = nn.Dropout(0.1)

```

```

8     def forward(self, x):
9         val, h_n = self.rnn(x)
10        val = self.dense(self.dropout(val[:, -1]))
11        return torch.flatten(val)
12    # Transfer to CUDA Device
13    model = NYSEModel().to(device)

```

Set the Hyperparameters and performance metric

```

1    # Hyperparameters and Performance Metrics
2    criterion = nn.MSELoss()
3    optimizer = optim.Adam(model.parameters(), lr=0.01)
4    Nepochs = 50

```

The training and validation section is shown below, it will plot the Accuracy for both on figure and update the figure on each epochs to see the progress of training.

```

1    # open Figure for Accuracies
2    fig, ax = plt.subplots()
3    ax.set_xlim(0, Nepochs)
4    ax.set_ylim(0, 1)
5    ax.set_xlabel('Epochs')
6    ax.set_ylabel('MSE')
7    ax.set_title('Training and Validation Loss')
8    train_acc_line, = ax.plot([], [], label='Training')
9    valid_acc_line, = ax.plot([], [], label='Validation')
10   plt.legend()
11   plt.draw()
12
13   # Training the Model
14   trainLoss, validLoss = [], []
15   for epoch in range(Nepochs):
16       model.train()
17       running_loss = 0
18       for input, target in train_loader:
19           optimizer.zero_grad()
20           # Add inputs and labels to device
21           input, target = input.to(device), target.to(device)
22           # Forward pass
23           predictions = model(input)
24           # Compute loss
25           loss = criterion(predictions, target)
26           # Backward pass
27           loss.backward()
28           optimizer.step()
29           # Compute loss
30           running_loss += loss.item()
31
32       trainLoss.append(running_loss / len(train_loader))
33       print(f'Epoch {epoch+1}/{Nepochs}, Loss: {running_loss / len(train_loader):.4f}')
34
35   model.eval()
36   running_loss = 0
37   with torch.no_grad():
38       for input, target in valid_loader:
39           # Add inputs and labels to device
40           input, target = input.to(device), target.to(device)
41           predictions = model(input)
42           # Compute loss
43           loss = criterion(predictions, target)
44           # Compute loss
45           running_loss += loss.item()

```

```

46     validLoss.append(running_loss/len(valid_loader))
47     # Update Plot
48     train_acc_line.set_data(range(1, len(trainLoss) + 1), trainLoss)
49     valid_acc_line.set_data(range(1, len(trainLoss) + 1), validLoss)
50     ax.relim()
51     plt.draw()
52     plt.pause(0.1)
53 plt.show()

```

Testing of the trained model is the essential part of the post-designing process, here we check how the model is performed on the test dataset

```

1 # Test Model
2 model.eval()
3 testLoss=[]
4 running_loss = 0
5 with torch.no_grad():
6     for input,target in test_loader:
7         # Add inputs and labels to device
8         input, target = input.to(device), target.to(device)
9         predictions = model(input)
10        # Compute loss
11        loss = criterion(predictions, target)
12        r2_metric.update(predictions, target)
13        mae_metric.update(predictions, target)
14        # Compute loss
15        running_loss += loss.item()
16    testLoss=running_loss/len(test_loader)
17 print('Testing Results')
18 print(f'R-Squared: {r2_metric.compute():0.4f}')
19 print(f'MSE: {testLoss:0.4f}')
20 print(f'MAE: {mae_metric.compute():0.4f}')

```

We will also forecast the trained model to see the output of the future volume of the stock.

```

1 # Forecasting Model
2 plt.plot(range(0,len(Yvalid)),Yvalid*stds[1]+means[1])
3 plt.plot(range(len(Yvalid)+1,len(Yvalid)+len(predictions.cpu().numpy())),
4          predictions.cpu().numpy()*stds[1]+means[1])
5 plt.xlabel('Time Step')
6 plt.ylabel('Log(Trading Volume)')
7 plt.legend(['Past Data','Forecasting'])
8 plt.show()

```

Hyperparameters and Performance Metrics

Hyperparameters

Optimizer : ADAM

Epochs : 50

Criterion : MSE Loss


Batch Size : 64

Performance Metric

R-Squared : 0.5298

MSE : 0.4838

MAE : 0.5257



Conclusions

We have fit the model and tried to get the forecasting based on the previous data. The R-squared values came to be 0.52 which is on the lower side, closer to value 1 treated as the good performance model and we have MSE of 0.4838.